**Congratulations! You passed!** Next Item

1. 1 / 1 point

Consider using this encoder-decoder model for machine translation.

This model is a "conditional language model" in the sense that the encoder portion (shown in green) is modeling the probability of the input sentence $x$.

- [ ] True
- [x] False

Correct

2. 1 / 1 point

In beam search, if you increase the beam width $B$, which of the following would you expect to be true? Check all that apply.

- [x] Beam search will run more slowly.

Correct

- [x] Beam search will use up more memory.

Correct

- [x] Beam search will generally find better solutions (i.e. do a better job maximizing $P(y \mid x)$)

Correct

- [ ] Beam search will converge after fewer steps.

Un-selected is correct

3. 1 / 1 point

In machine translation, if we carry out beam search without using sentence normalization, the algorithm will tend to output overly short translations.

- [x] True

Correct

- [ ] False

4. 1 / 1 point

Suppose you are building a speech recognition system, which uses an RNN model to map from audio clip $x$ to a text transcript $y$. Your algorithm uses beam search to try to find the value of $y$ that maximizes $P(y \mid x)$.

On a dev set example, given an input audio clip, your algorithm outputs the transcript $\hat{y}$ = "I'm building an A Eye system in Silly con Valley.", whereas a human gives a much superior transcript $y^*$ = "I'm building an AI system in Silicon Valley."

According to your model,

$P(\hat{y} \mid x) = 1.09 * 10^{-}7$

$P(y^* \mid x) = 7.21 * 10^{-}8$

Would you expect increasing the beam width B to help correct this example?

- [x] No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.

Correct

- [ ] No, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.
- [ ] Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the RNN rather than to the search algorithm.
- [ ] Yes, because $P(y^* \mid x) \leq P(\hat{y} \mid x)$ indicates the error should be attributed to the search algorithm rather than to the RNN.

5. 1 / 1 point

Continuing the example from Q4, suppose you work on your algorithm for a few more weeks, and now find that for the vast majority of examples on which your algorithm makes a mistake, $P(y^* \mid x) > P(\hat{y} \mid x)$. This suggest you should focus your attention on improving the search algorithm.

- [x] True.

Correct

- [ ] False.

6. 1 / 1 point

Consider the attention model for machine translation.

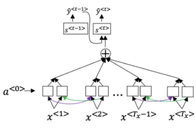

Further, here is the formula for $\alpha^{<t,t'>}$.

$$\alpha^{<t,t'>} = \frac{\exp(e^{<t,t'>})}{\sum_{t'=1}^{T_x} \exp(e^{<t,t'>})}$$

Which of the following statements about $\alpha^{<t,t'>}$ are true? Check all that apply.

- [x] We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t'>}$ that are highly relevant to the value the network should output for $y^{<t>}$. (Note the indices in the superscripts.)

Correct

- [ ] We expect $\alpha^{<t,t'>}$ to be generally larger for values of $a^{<t>}$ that are highly relevant to the value the network should output for $y^{<t'>}$. (Note the indices in the superscripts.)

Un-selected is correct

- [ ] $\sum_t \alpha^{<t,t'>} = 1$ (Note the summation is over $t$.)

Un-selected is correct

- [x] $\sum_{t'} \alpha^{<t,t'>} = 1$ (Note the summation is over $t'$.)

Correct

7. 1 / 1 point

The network learns where to "pay attention" by learning the values $e^{<t,t'>}$, which are computed using a small neural network:

We can't replace $s^{<t-1>}$ with $s^{<t>}$ as an input to this neural network. This is because $s^{<t>}$ depends on $\alpha^{<t,t'>}$ which in turn depends on $e^{<t,t'>}$; so at the time we need to evalute this network, we haven't computed $s^{<t>}$ yet.

- [x] True

Correct

- [ ] False

8. 1 / 1 point

Compared to the encoder-decoder model shown in Question 1 of this quiz (which does not use an attention mechanism), we expect the attention model to have the greatest advantage when:

- [x] The input sequence length $T_x$ is large.

Correct

- [ ] The input sequence length $T_x$ is small.

9. 1 / 1 point

Under the CTC model, identical repeated characters not separated by the "blank" character (_) are collapsed. Under the CTC model, what does the following string collapse to?

__c_oo_o_kk___b_ooooo__oo_kkk

- [ ] cokbok
- [x] cookbook

Correct

- [ ] cook book
- [ ] coookkboooooookkk

10. 1 / 1 point

In trigger word detection, $x^{<t>}$ is:

- [x] Features of the audio (such as spectrogram features) at time $t$.

Correct

- [ ] The $t$-th input word, represented as either a one-hot vector or a word embedding.
- [ ] Whether the trigger word is being said at time $t$.
- [ ] Whether someone has just finished saying the trigger word at time $t$.